

Lab - 2 - Data Exploration

Step 1. Import the necessary libraries

```
import pandas as pd
```

Step 2. Import the dataset from this [address](#).

Step 3. Assign it to a variable called users and use the 'user_id' as index

```
df = pd.read_csv("users.txt", sep="|", index_col=[0])
```

Step 4. See the first 25 entries

```
df.head(25)
```

	age	gender	occupation	zip_code
user_id				
1	24	M	technician	85711
2	53	F	other	94043
3	23	M	writer	32067
4	24	M	technician	43537
5	33	F	other	15213
6	42	M	executive	98101
7	57	M	administrator	91344
8	36	M	administrator	05201
9	29	M	student	01002
10	53	M	lawyer	90703
11	39	F	other	30329
12	28	F	other	06405
13	47	M	educator	29206
14	45	M	scientist	55106
15	49	F	educator	97301
16	21	M	entertainment	10309
17	30	M	programmer	06355
18	35	F	other	37212
19	40	M	librarian	02138
20	42	F	homemaker	95660
21	26	M	writer	30068
22	25	M	writer	40206
23	30	F	artist	48197
24	21	F	artist	94533
25	39	M	engineer	55107

Step 5. See the last 10 entries

```
df.tail(10)
```

	age	gender	occupation	zip_code
user_id				
934	61	M	engineer	22902
935	42	M	doctor	66221
936	24	M	other	32789
937	48	M	educator	98072
938	38	F	technician	55038
939	26	F	student	33319
940	32	M	administrator	02215
941	20	M	student	97229
942	48	F	librarian	78209
943	22	M	student	77841

Step 6. What is the number of observations in the dataset?

```
#0 indicates row
```

```
df.shape[0]
```

```
943
```

Step 7. What is the number of columns in the dataset?

```
#1 indicates column
```

```
df.shape[1]
```

```
4
```

Step 8. Print the name of all the columns.

```
df.columns
```

```
Index(['age', 'gender', 'occupation', 'zip_code'], dtype='object')
```

Step 9. How is the dataset indexed?

```
# "the index" (aka "the labels")
```

```
df.index
```

```
Int64Index([ 1,  2,  3,  4,  5,  6,  7,  8,  9, 10,
             ...,
            934, 935, 936, 937, 938, 939, 940, 941, 942, 943],
            dtype='int64', name='user_id', length=943)
```

Step 10. What is the data type of each column?

```
df.dtypes
```

```
age          int64
gender       object
occupation   object
zip_code     object
dtype: object
```

Step 11. Print only the occupation column

```
df["occupation"]

user_id
1      technician
2           other
3          writer
4      technician
5           other
...
939         student
940  administrator
941         student
942        librarian
943         student
Name: occupation, Length: 943, dtype: object
```

Step 12. How many different occupations are in this dataset?

```
#print number of unique occupation

print(df.occupation.nunique())

#print all unique occupation in tuple

print(df.occupation.unique())

21
['technician' 'other' 'writer' 'executive' 'administrator' 'student'
 'lawyer' 'educator' 'scientist' 'entertainment' 'programmer'
 'librarian'
 'homemaker' 'artist' 'engineer' 'marketing' 'none' 'healthcare'
 'retired'
 'salesman' 'doctor']
```

Step 13. What is the most frequent occupation?

```
'student'
```

Step 14. Summarize the DataFrame.

```
df.describe()
```

	age
count	943.000000
mean	34.051962
std	12.192740
min	7.000000
25%	25.000000
50%	31.000000
75%	43.000000
max	73.000000

Step 15. Summarize all the columns

```
df.describe(include="all")
```

	age	gender	occupation	zip_code
count	943.000000	943	943	943
unique	NaN	2	21	795
top	NaN	M	student	55414
freq	NaN	670	196	9
mean	34.051962	NaN	NaN	NaN
std	12.192740	NaN	NaN	NaN
min	7.000000	NaN	NaN	NaN
25%	25.000000	NaN	NaN	NaN
50%	31.000000	NaN	NaN	NaN
75%	43.000000	NaN	NaN	NaN
max	73.000000	NaN	NaN	NaN

Step 16. Summarize only the occupation column

```
df.occupation.describe()
```

count	943
unique	21
top	student
freq	196

Name: occupation, dtype: object

Step 17. What is the mean age of users?

```
round(df.age.mean())
```

34

Step 18. What is the age with least occurrence?

```
df.age.value_counts().tail()
```

```
7      1
66     1
11     1
10     1
73     1
Name: age, dtype: int64
```

Lab - 3 - Data Exploration

1) First, you need to read the titanic dataset from local disk and display first five records

```
import pandas as pd
```

```
df = pd.read_csv("titanic.csv")
```

```
df.head(5)
```

	PassengerId	Survived	Pclass	\
0	1	0	3	
1	2	1	1	
2	3	1	3	
3	4	1	1	
4	5	0	3	

	SibSp	\	Name	Sex	Age
0			Braund, Mr. Owen Harris	male	22.0
1					
1	1		Cumings, Mrs. John Bradley (Florence Briggs Th...	female	38.0
1					
2			Heikkinen, Miss. Laina	female	26.0
0					
3			Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0
1					
4			Allen, Mr. William Henry	male	35.0
0					

	Parch		Ticket	Fare	Cabin	Embarked
0	0		A/5 21171	7.2500	NaN	S
1	0		PC 17599	71.2833	C85	C
2	0	STON/O2.	3101282	7.9250	NaN	S
3	0		113803	53.1000	C123	S
4	0		373450	8.0500	NaN	S

2) Identify Nominal, Ordinal, Binary and Numeric attributes from data sets and display all values.

```
print("Nominal ")
print(df["Name"])
print(df["Ticket"])
print(df["Embarked"])
print(df["Cabin"])
print("Ordinal")
```

```

print(df["Pclass"])
print("Binary")
print(df["Sex"])
print(df["Survived"])
print("Numeric")
print(df["SibSp"])
print(df["PassengerId"])
print(df["Fare"])
print(df["Age"])
print(df["Parch"])

```

Nominal

```

0          Braund, Mr. Owen Harris
1  Cumings, Mrs. John Bradley (Florence Briggs Th...
2          Heikkinen, Miss. Laina
3  Futrelle, Mrs. Jacques Heath (Lily May Peel)
4          Allen, Mr. William Henry
...
886          Montvila, Rev. Juozas
887          Graham, Miss. Margaret Edith
888  Johnston, Miss. Catherine Helen "Carrie"
889          Behr, Mr. Karl Howell
890          Dooley, Mr. Patrick

```

Name: Name, Length: 891, dtype: object

```

0      A/5 21171
1      PC 17599
2  STON/O2. 3101282
3      113803
4      373450
...
886      211536
887      112053
888  W./C. 6607
889      111369
890      370376

```

Name: Ticket, Length: 891, dtype: object

```

0      S
1      C
2      S
3      S
4      S
...
886      S
887      S
888      S
889      C
890      Q

```

Name: Embarked, Length: 891, dtype: object

```

0      NaN
1      C85

```

```

2      NaN
3      C123
4      NaN
...
886    NaN
887    B42
888    NaN
889    C148
890    NaN
Name: Cabin, Length: 891, dtype: object
Ordinal
0      3
1      1
2      3
3      1
4      3
..
886    2
887    1
888    3
889    1
890    3
Name: Pclass, Length: 891, dtype: int64
Binary
0      male
1      female
2      female
3      female
4      male
...
886    male
887    female
888    female
889    male
890    male
Name: Sex, Length: 891, dtype: object
0      0
1      1
2      1
3      1
4      0
..
886    0
887    1
888    0
889    1
890    0
Name: Survived, Length: 891, dtype: int64
Numeric

```



```

0      1
1      1
2      0
3      1
4      0
...
886    0
887    0
888    1
889    0
890    0
Name: SibSp, Length: 891, dtype: int64
0      1
1      2
2      3
3      4
4      5
...
886    887
887    888
888    889
889    890
890    891
Name: PassengerId, Length: 891, dtype: int64
0      7.2500
1     71.2833
2      7.9250
3     53.1000
4      8.0500
...
886    13.0000
887    30.0000
888    23.4500
889    30.0000
890     7.7500
Name: Fare, Length: 891, dtype: float64
0     22.0
1     38.0
2     26.0
3     35.0
4     35.0
...
886    27.0
887    19.0
888     NaN
889    26.0
890    32.0
Name: Age, Length: 891, dtype: float64
0      0

```

```

1      0
2      0
3      0
4      0
...
886    0
887    0
888    2
889    0
890    0
Name: Parch, Length: 891, dtype: int64

```

3) Identify symmetric and asymmetric binary attributes from data sets and display all values.

```

print("Symmetric",df["Sex"])
print("Asymmetric",df["Survived"])

Symmetric 0      male
1      female
2      female
3      female
4      male
...
886      male
887      female
888      female
889      male
890      male
Name: Sex, Length: 891, dtype: object
Asymmetric 0      0
1      1
2      1
3      1
4      0
...
886    0
887    1
888    0
889    1
890    0
Name: Survived, Length: 891, dtype: int64

```

4) For each quantitative attribute, calculate its average, standard deviation, minimum, mode, range and maximum values.

```

from pandas.api.types import is_numeric_dtype
for column in df.columns:
    if(is_numeric_dtype(df[column])):

```

```
print(column,":")
print("\tMean : ",df[column].mean())
print("\tStandard Deviation : ",df[column].std())
print("\tMinimum : ",df[column].min())
print("\tRange : ",df[column].max()-df[column].min())
print("\tMax : ",df[column].max())
if column!="PassengerId":
    print("\tMode : ",df[column].mode()[0])
```

PassengerId :

Mean : 446.0
Standard Deviation : 257.3538420152301
Minimum : 1
Range : 890
Max : 891

Survived :

Mean : 0.3838383838383838
Standard Deviation : 0.4865924542648585
Minimum : 0
Range : 1
Max : 1
Mode : 0

Pclass :

Mean : 2.308641975308642
Standard Deviation : 0.8360712409770513
Minimum : 1
Range : 2
Max : 3
Mode : 3

Age :

Mean : 29.69911764705882
Standard Deviation : 14.526497332334044
Minimum : 0.42
Range : 79.58
Max : 80.0
Mode : 24.0

SibSp :

Mean : 0.5230078563411896
Standard Deviation : 1.1027434322934275
Minimum : 0
Range : 8
Max : 8
Mode : 0

Parch :

Mean : 0.38159371492704824
Standard Deviation : 0.8060572211299559
Minimum : 0
Range : 6
Max : 6
Mode : 0

```
Fare :
  Mean : 32.2042079685746
  Standard Deviation : 49.693428597180905
  Minimum : 0.0
  Range : 512.3292
  Max : 512.3292
  Mode : 8.05
```

6) For the qualitative attribute (class), count the frequency for each of its distinct values.

```
df["Pclass"].value_counts()
3      491
1      216
2      184
Name: Pclass, dtype: int64
```

7) It is also possible to display the summary for all the attributes simultaneously in a table using the describe() function. If an attribute is quantitative, it will display its mean, standard deviation and various quantiles (including minimum, median, and maximum) values. If an attribute is qualitative, it will display its number of unique values and the top (most frequent) values.

```
df.describe(include="all")
```

	PassengerId	Survived	Pclass	Name
Sex \				
count	891.000000	891.000000	891.000000	891
unique	NaN	NaN	NaN	891
top	NaN	NaN	NaN	Braund, Mr. Owen Harris
male				
freq	NaN	NaN	NaN	1
577				
mean	446.000000	0.383838	2.308642	NaN
NaN				
std	257.353842	0.486592	0.836071	NaN
NaN				
min	1.000000	0.000000	1.000000	NaN
NaN				
25%	223.500000	0.000000	2.000000	NaN
NaN				
50%	446.000000	0.000000	3.000000	NaN
NaN				

75%	668.500000	1.000000	3.000000			NaN
NaN						
max	891.000000	1.000000	3.000000			NaN
NaN						
	Age	SibSp	Parch	Ticket	Fare	
Cabin \						
count	714.000000	891.000000	891.000000	891	891.000000	
204						
unique	NaN	NaN	NaN	681	NaN	
147						
top	NaN	NaN	NaN	347082	NaN	B96
B98						
freq	NaN	NaN	NaN	7	NaN	
4						
mean	29.699118	0.523008	0.381594	NaN	32.204208	
NaN						
std	14.526497	1.102743	0.806057	NaN	49.693429	
NaN						
min	0.420000	0.000000	0.000000	NaN	0.000000	
NaN						
25%	20.125000	0.000000	0.000000	NaN	7.910400	
NaN						
50%	28.000000	0.000000	0.000000	NaN	14.454200	
NaN						
75%	38.000000	1.000000	0.000000	NaN	31.000000	
NaN						
max	80.000000	8.000000	6.000000	NaN	512.329200	
NaN						
	Embarked					
count	889					
unique	3					
top	S					
freq	644					
mean	NaN					
std	NaN					
min	NaN					
25%	NaN					
50%	NaN					
75%	NaN					
max	NaN					

8) For multivariate statistics, you can compute the covariance and correlation between pairs of attributes.

```
df.cov()
```


	PassengerId	Survived	Pclass	Age	SibSp
\ PassengerId	66231.000000	-0.626966	-7.561798	138.696504	-16.325843
Survived	-0.626966	0.236772	-0.137703	-0.551296	-0.018954
Pclass	-7.561798	-0.137703	0.699015	-4.496004	0.076599
Age	138.696504	-0.551296	-4.496004	211.019125	-4.163334
SibSp	-16.325843	-0.018954	0.076599	-4.163334	1.216043
Parch	-0.342697	0.032017	0.012429	-2.344191	0.368739
Fare	161.883369	6.221787	-22.830196	73.849030	8.748734
	Parch	Fare			
PassengerId	-0.342697	161.883369			
Survived	0.032017	6.221787			
Pclass	0.012429	-22.830196			
Age	-2.344191	73.849030			
SibSp	0.368739	8.748734			
Parch	0.649728	8.661052			
Fare	8.661052	2469.436846			
df.corr()					
	PassengerId	Survived	Pclass	Age	SibSp
Parch \					
PassengerId	1.000000	-0.005007	-0.035144	0.036847	-0.057527
Survived	-0.005007	1.000000	-0.338481	-0.077221	-0.035322
Pclass	-0.035144	-0.338481	1.000000	-0.369226	0.083081
Age	0.036847	-0.077221	-0.369226	1.000000	-0.308247
SibSp	-0.057527	-0.035322	0.083081	-0.308247	1.000000
Parch	-0.001652	0.081629	0.018443	-0.189119	0.414838
Fare	0.012658	0.257307	-0.549500	0.096067	0.159651
	Fare				
PassengerId	0.012658				
Survived	0.257307				
Pclass	-0.549500				
Age	0.096067				
SibSp	0.159651				

```
Parch      0.216225  
Fare       1.000000
```

9) Display the histogram for Age attribute by discretizing it into 8 separate bins and counting the frequency for each bin.

```
df["Age"].hist(bins=8)
```


```
<AxesSubplot:>
```



10) A boxplot can also be used to show the distribution of values for each attribute.

```
df.boxplot(column=["Age", "Fare"])
```



```
<AxesSubplot:>
```



11) Display scatter plot for any 5 pair of attributes , we can use a scatter plot to visualize their joint distribution.

```
df.plot.scatter(x="Age",y="Fare")
```

```
<AxesSubplot:xlabel='Age', ylabel='Fare'>
```

Lab - 4 - Data Preprocessing

1) First, you need to read the titanic dataset from local disk and display Last five records

```
import pandas as pd
```

```
df = pd.read_csv("titanic.csv")
```

```
df.tail(5)
```

	PassengerId	Survived	Pclass	Name \
886	887	0	2	Montvila, Rev. Juozas
887	888	1	1	Graham, Miss. Margaret Edith
888	889	0	3	Johnston, Miss. Catherine Helen "Carrie"
889	890	1	1	Behr, Mr. Karl Howell
890	891	0	3	Dooley, Mr. Patrick

	Sex	Age	SibSp	Parch	Ticket	Fare	Cabin	Embarked
886	male	27.0	0	0	211536	13.00	NaN	S
887	female	19.0	0	0	112053	30.00	B42	S
888	female	NaN	1	2	W./C. 6607	23.45	NaN	S
889	male	26.0	0	0	111369	30.00	C148	C
890	male	32.0	0	0	370376	7.75	NaN	Q

2) Handle Missing Values in data set [use dropna(), fillna(), and interpolate]

```
df.isnull().sum()
```

PassengerId	0
Survived	0
Pclass	0
Name	0
Sex	0
Age	177
SibSp	0
Parch	0
Ticket	0
Fare	0
Cabin	687

```
Embarked      2
dtype: int64

df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 891 entries, 0 to 890
Data columns (total 12 columns):
#   Column      Non-Null Count  Dtype
---  -
0   PassengerId  891 non-null    int64
1   Survived     891 non-null    int64
2   Pclass       891 non-null    int64
3   Name         891 non-null    object
4   Sex          891 non-null    object
5   Age         714 non-null    float64
6   SibSp        891 non-null    int64
7   Parch        891 non-null    int64
8   Ticket       891 non-null    object
9   Fare         891 non-null    float64
10  Cabin        204 non-null    object
11  Embarked     889 non-null    object
dtypes: float64(2), int64(5), object(5)
memory usage: 83.7+ KB

newdf =
df.fillna({'Age':df['Age'].mean(),'Cabin':"Y28",'Embarked':"Y"})

newdf

   PassengerId  Survived  Pclass  \
0             1         0       3
1             2         1       1
2             3         1       3
3             4         1       1
4             5         0       3
..          ...       ...       ...
886          887         0       2
887          888         1       1
888          889         0       3
889          890         1       1
890          891         0       3

   Name      Sex
Age  \
0      Braund, Mr. Owen Harris  male
22.000000
1  Cumings, Mrs. John Bradley (Florence Briggs Th...  female
38.000000
2      Heikkinen, Miss. Laina  female
```

```

26.000000
3      Futrelle, Mrs. Jacques Heath (Lily May Peel)  female
35.000000
4      Allen, Mr. William Henry  male
35.000000
...
...
886      Montvila, Rev. Juozas  male
27.000000
887      Graham, Miss. Margaret Edith  female
19.000000
888      Johnston, Miss. Catherine Helen "Carrie"  female
29.699118
889      Behr, Mr. Karl Howell  male
26.000000
890      Dooley, Mr. Patrick  male
32.000000

```

	SibSp	Parch	Ticket	Fare	Cabin	Embarked
0	1	0	A/5 21171	7.2500	Y28	S
1	1	0	PC 17599	71.2833	C85	C
2	0	0	STON/02. 3101282	7.9250	Y28	S
3	1	0	113803	53.1000	C123	S
4	0	0	373450	8.0500	Y28	S
...
886	0	0	211536	13.0000	Y28	S
887	0	0	112053	30.0000	B42	S
888	1	2	W./C. 6607	23.4500	Y28	S
889	0	0	111369	30.0000	C148	C
890	0	0	370376	7.7500	Y28	Q

```
[891 rows x 12 columns]
```

```
newdf.isnull().sum()
```

```

PassengerId    0
Survived        0
Pclass          0
Name            0
Sex             0
Age            0
SibSp           0
Parch           0
Ticket          0
Fare            0
Cabin           0
Embarked        0
dtype: int64

```

```
for i in df['Age']:
    if i<1:
        print(i)
```

```
0.83
0.92
0.75
0.75
0.67
0.42
0.83
```

3) Apply Scaling to AGE attribute with min max, decimal scaling and z score.

```
newdf['newAge'] = (newdf['Age'] -
newdf['Age'].min())/(newdf['Age'].max() - newdf['Age'].min())
```

```
newdf.describe()
```

	PassengerId	Survived	Pclass	Age	SibSp	\
count	891.000000	891.000000	891.000000	891.000000	891.000000	
mean	446.000000	0.383838	2.308642	29.699118	0.523008	
std	257.353842	0.486592	0.836071	13.002015	1.102743	
min	1.000000	0.000000	1.000000	0.420000	0.000000	
25%	223.500000	0.000000	2.000000	22.000000	0.000000	
50%	446.000000	0.000000	3.000000	29.699118	0.000000	
75%	668.500000	1.000000	3.000000	35.000000	1.000000	
max	891.000000	1.000000	3.000000	80.000000	8.000000	

	Parch	Fare	newAge
count	891.000000	891.000000	891.000000
mean	0.381594	32.204208	0.367921
std	0.806057	49.693429	0.163383
min	0.000000	0.000000	0.000000
25%	0.000000	7.910400	0.271174
50%	0.000000	14.454200	0.367921
75%	0.000000	31.000000	0.434531
max	6.000000	512.329200	1.000000

```
newdf = df
newdf['newAge'] = df['Age']/10**len(str(int(df['Age'].max())))
```

```
df
```

	PassengerId	Survived	Pclass	\
0	1	0	3	
1	2	1	1	
2	3	1	3	
3	4	1	1	

```

4          5          0          3
..      ...      ...      ...
886      887          0          2
887      888          1          1
888      889          0          3
889      890          1          1
890      891          0          3

```

```

                                     Name      Sex   Age
SibSp \
0                                     Braund, Mr. Owen Harris    male  22.0
1
1      Cumings, Mrs. John Bradley (Florence Briggs Th...  female  38.0
1
2                                     Heikkinen, Miss. Laina  female  26.0
0
3      Futrelle, Mrs. Jacques Heath (Lily May Peel)  female  35.0
1
4                                     Allen, Mr. William Henry    male  35.0
0
..                                     ...      ...   ...
...
886                                     Montvila, Rev. Juozas    male  27.0
0
887                                     Graham, Miss. Margaret Edith  female  19.0
0
888      Johnston, Miss. Catherine Helen "Carrie"  female   NaN
1
889                                     Behr, Mr. Karl Howell    male  26.0
0
890                                     Dooley, Mr. Patrick    male  32.0
0

```

```

      Parch      Ticket      Fare Cabin Embarked  newAge
0          0      A/5 21171   7.2500   NaN      S    0.22
1          0      PC 17599  71.2833   C85      C    0.38
2          0  STON/O2. 3101282   7.9250   NaN      S    0.26
3          0      113803  53.1000  C123      S    0.35
4          0      373450   8.0500   NaN      S    0.35
..      ...      ...      ...   ...      ...   ...
886          0      211536  13.0000   NaN      S    0.27
887          0      112053  30.0000   B42      S    0.19
888          2      W./C. 6607  23.4500   NaN      S    NaN
889          0      111369  30.0000  C148      C    0.26
890          0      370376   7.7500   NaN      Q    0.32

```

```
[891 rows x 13 columns]
```

```
df.corr()
```

	PassengerId	Survived	Pclass	Age	SibSp
Parch \ PassengerId	1.000000	-0.005007	-0.035144	0.036847	-0.057527
0.001652					
Survived	-0.005007	1.000000	-0.338481	-0.077221	-0.035322
0.081629					
Pclass	-0.035144	-0.338481	1.000000	-0.369226	0.083081
0.018443					
Age	0.036847	-0.077221	-0.369226	1.000000	-0.308247
0.189119					
SibSp	-0.057527	-0.035322	0.083081	-0.308247	1.000000
0.414838					
Parch	-0.001652	0.081629	0.018443	-0.189119	0.414838
1.000000					
Fare	0.012658	0.257307	-0.549500	0.096067	0.159651
0.216225					
newAge	0.036847	-0.077221	-0.369226	1.000000	-0.308247
0.189119					

	Fare	newAge
PassengerId	0.012658	0.036847
Survived	0.257307	-0.077221
Pclass	-0.549500	-0.369226
Age	0.096067	1.000000
SibSp	0.159651	-0.308247
Parch	0.216225	-0.189119
Fare	1.000000	0.096067
newAge	0.096067	1.000000

```
df[['Age', 'newAge']].corr()
```

	Age	newAge
Age	1.0	1.0
newAge	1.0	1.0